

Reinforcement Learning

November 12, 2019

1 IE 534 HW: Reinforcement Learning

v1, Designed by Yuanyi Zhong, 2019

In this assignment, we will experiment with the (deep) reinforcement learning algorithms covered in the lecture. In particular, you will implement variants of the popular DQN (Deep Q-Network) (1) and A2C (Advantage Actor-Critic) (2) algorithms (by the same first author! orz), and test your implementation on both a small example (CartPole problem) and an Atari game (Breakout game). We focus on model-free algorithms rather than model-based ones, because neural nets are easier applicable and more popular nowadays in the model-free setting. (When the system dynamic is known or can be easily inferred, model-based can sometimes do better.)

The assignment breaks into **three parts**:

- **In Part I** (50 pts), you basically need to follow the instructions in this notebook to do a little bit of coding. We'll be able to see if your code trains by testing against the CartPole environment provided by the OpenAI gym package. We'll generate some plots that are required for grading.
- **In Part II** (40 pts), you'll copy your code onto Blue Waters (or actually any good server.), and run a much larger-scale experiment with the Breakout game. Hopefully, you can teach the computer to play Breakout in less than half a day! Share your final game score in this notebook. **This part will take at least a day. Please start early!!**
- **In Part III** (10 pts), you'll be asked to think about a few questions. These questions are mostly open-ended. Please write down your thoughts on them.

Finally, after you finished everything in this notebook (**code snippets C1-C5, plots P1-P5, question answers Q1-Q5**), please save the notebook, and export to a PDF (or an HTML file), and submit:

1. the **.ipynb notebook and exported .pdf/.html file**, PDF is preferred (I usually do File -> Print Preview -> use Chrome's Save as PDF);
2. your code (**Algo.py, Model.py files**);
3. job artifacts (**.log files** only, pytorch models and images not required) to Compass 2g for grading.

PS: Remember to save your notebook occasionally as you work through it!

References

- (1) Mnih, V., Kavukcuoglu, K., Silver, D., Rusu, A.A., Veness, J., Bellemare, M.G., Graves, A.,
- (2) Mnih, V., Badia, A.P., Mirza, M., Graves, A., Lillicrap, T., Harley, T., Silver, D. and Kav
- (3) A useful tutorial: <https://spinningup.openai.com/en/latest/>
- (4) *Useful code references*: <https://github.com/deepmind/bsuite>; <https://github.com/openai/base>

First of all, **enter your NetID here** in the cell below: Your NetID: yuanyiz2

1.1 Part I: DQN and A2C on CartPole

This part is designed to run on your own local laptop/PC.

Before you start, there are some python dependencies: `pytorch`, `gym`, `numpy`, `multiprocessing`, `matplotlib`. Please install them correctly. You can install `pytorch` following instruction here <https://pytorch.org/get-started/locally/>. The code is compatible with PyTorch 0.4.x ~ 1.x. PyTorch 1.1 with cuda 10.0 worked for me (`conda install pytorch==1.1.0 torchvision==0.3.0 cudatoolkit=10.0 -c pytorch`).

Please **always** run the code cell below each time you open this notebook, to make sure `gym` is installed and to enable `autoreload` which **allows code changes to be effective immediately**. So if you changed `Algo.py` or `Model.py` but the test codes are not reflecting your changes, restart the notebook kernel and run this cell!!

```
In [1]: # install openai gym
        %pip install gym
        # enable autoreload
        %load_ext autoreload
        %autoreload 2
```

```
Requirement already satisfied: gym in c:\users\bill\anaconda3\lib\site-packages (0.14.0)
Requirement already satisfied: scipy in c:\users\bill\anaconda3\lib\site-packages (from gym) (1.
Requirement already satisfied: cloudpickle~=1.2.0 in c:\users\bill\anaconda3\lib\site-packages (
Requirement already satisfied: pygame<=1.3.2,>=1.2.0 in c:\users\bill\anaconda3\lib\site-package
Requirement already satisfied: numpy>=1.10.4 in c:\users\bill\anaconda3\lib\site-packages (from
Requirement already satisfied: six in c:\users\bill\anaconda3\lib\site-packages (from gym) (1.12
Requirement already satisfied: future in c:\users\bill\anaconda3\lib\site-packages (from pygame<
Note: you may need to restart the kernel to use updated packages.
```

1.1.1 1.1 Code Structure

The code is structured in 5 python files:

- `Main.py`: contains the main entry point and training loop
- `Model.py`: constructs the torch neural network modules

- `Env.py`: contains the environment simulations interface, based on openai gym
- `Algo.py`: implements the DQN and A2C algorithms
- `Replay.py`: implements the experience replay buffer for DQN
- `Draw.py`: saves some game snapshots to jpeg files

Some parts of the code `Model.py` and `Algo.py` are left blank for you to complete. You are not required to modify the other parts (unless, of course, you want to boost the performance!). This is kind of a minimalist implementation, and might be different from the other code on the internet in details. You're welcomed to improve it, after you've finished all the required things of this assignment.

1.1.2 1.2 OpenAI gym and CartPole environment

OpenAI developed python package gym a while ago to facilitate RL research. gym provides a common interface between the program and the environments. For instance, the code cell below will create the CartPole environment. A window will show up when you run the code. The goal is to keep adjusting the cart so that the pole stays in its upright position.

A demo video from OpenAI:

gym also provides interface to Atari games. However, installing package `atari-py` is not easy on Windows/Mac, so we won't demonstrate it here. More info: <http://gym.openai.com/docs/>.

```
In [2]: import time
import gym
env = gym.make('CartPole-v1')
env.reset()
for _ in range(200):
    env.render()
    state, reward, done, _ = env.step(env.action_space.sample()) # take a random action
    if done: break
    time.sleep(0.15)
env.close()
```

1.1.3 1.3 Deep Q Learning

A little recap on DQN. We learned from lecture that Q-Learning is a model-free reinforcement learning algorithm. It falls into the off-policy type algorithm since it can utilize past experiences stored in a buffer. It also falls into the (approximate) dynamic programming type algorithm, since it tries to learn an optimal state-action value function using time difference (TD) errors. Q Learning is particularly interesting because it exploits the optimality structure in MDP. It's related to the Hamilton-Jacobi-Bellman equation in classical control.

For MDP

$$M = (S, A, P, r, \gamma)$$

where S is the state space, A is the action space, P is the transition dynamic, $r(s, a)$ is a reward function $S \times A \mapsto R$, and γ is the discount factor.

The tabular case (when S, A are finite), Q-Learning does the following value iteration update repeatedly when collecting experience (s_t, a_t, r_t) (η is the learning rate):

$$Q^{new}(s_t, a_t) \leftarrow Q^{old}(s_t, a_t) + \eta \left(r_t + \gamma \max_{a' \in A} Q^{old}(s_t, a') - Q^{old}(s_t, a_t) \right).$$

With function approximation, meaning model $Q(s, a)$ with a function $Q_\theta(s, a)$ parameterized by θ , we arrive at the Fitted Q Iteration (FQI) algorithm, or better known as Deep Q Learning if the function class is neural networks. Q-Learning with neural network as function approximator was known long ago, but it was only recently (year 2013) that DeepMind made this algorithm actually work on Atari games. Deep Q Learning iteratively optimize the following objective:

$$\theta_{new} \leftarrow \arg \min_{\theta} \mathbb{E}_{(s,a,r,s') \sim D} \left(r + \gamma \max_{a' \in A} Q_{\theta_{old}}(s', a') - Q_{\theta}(s, a) \right)^2.$$

Therefore, with a batch of $\{(s^i, a^i, r^i, s'^i)\}_{i=1}^N$ sampled from the replay buffer, we can build a loss function L in pytorch:

$$L(\theta) = \frac{1}{N} \sum_{i=1}^N \left(r^i + \gamma \max_{a' \in A} Q_{\theta_{old}}(s'^i, a') - Q_{\theta}(s^i, a^i) \right)^2,$$

and run the usual gradient descent on θ with a pytorch optimizer.

Exploration Exploration, as the word suggests, refers to explore novel unvisited states in RL. The FQI (or DQN) needs an exploratory datasets to work well. The common way to produce exploratory dataset is through randomization, such as the ϵ -greedy exploration strategy we will implement in this assignment. - ϵ -greedy exploration:

At training iteration it , the agent chooses to play

$$a = \begin{cases} \arg \max_a Q_{\theta}(s, a) & \text{with probability } 1 - \epsilon_{it} , \\ \text{a random action } a \in A & \text{with probability } \epsilon_{it} . \end{cases}$$

And ϵ_{it} is annealed, for example, linearly from 1 to 0.01 as training progresses until iteration it_{decay} :

$$\epsilon_{it} = \max \left\{ 0.01, 1 + (0.01 - 1) \frac{it}{it_{decay}} \right\}.$$

Two Caveats

1. There's an improvement on DQN called Double-DQN with the following loss L , which has shown to be empirically more stable than the original DQN loss described above. You may want to implement the improved one in your code:

$$L(\theta) = \frac{1}{N} \sum_{i=1}^N \left(r^i + \gamma Q_{\theta_{old}}(s'^i, \arg \max_{a' \in A} Q_{\theta}(s'^i, a')) - Q_{\theta}(s^i, a^i) \right)^2.$$

2. Huber loss (a.k.a smooth L1 loss) is commonly used to reduce the effect of extreme values:

$$L(\theta) = \frac{1}{N} \sum_{i=1}^N \text{Huber} \left(r^i + \gamma Q_{\theta_{old}}(s'^i, \arg \max_{a' \in A} Q_{\theta}(s'^i, a')) - Q_{\theta}(s^i, a^i) \right)$$

You can look up the pytorch document here: <https://pytorch.org/docs/stable/nn.functional.html#smooth-l1-loss>

C1 (5 pts): Complete the code for the two layered fully connected network class `TwoLayerFCNet` in file `Model.py` And run the cell below to test the output shape of your module.

```
In [3]: ## Test code
        from Model import TwoLayerFCNet
        import torch
        net = TwoLayerFCNet(n_in=4, n_hidden=16, n_out=5)
        x = torch.randn(10, 4)
        y = net(x)
        assert y.shape == (10, 5), "ERROR: network output has the wrong shape!"
        print ("Output shape test passed!")
```

Output shape test passed!

C2 (5 pts): Complete the code for ϵ -greedy exploration strategy in function `DQN.act` in file `'Algo.py'` And run the cell below to test it.

```
In [4]: ## Test code
        from Algo import DQN
        class Nothing: pass
        dummy = Nothing()
        dummy.eps_decay = 500000

        dummy.num_act_steps = 0
        eps = DQN.compute_epsilon(dummy)
        assert abs( eps - 1.0 ) < 0.01, "ERROR: compute_epsilon at t=0 should be 1 but got %f!"

        dummy.num_act_steps = 250000
        eps = DQN.compute_epsilon(dummy)
        assert abs( eps - 0.505 ) < 0.01, "ERROR: compute_epsilon at t=250000 should around .505"

        dummy.num_act_steps = 500000
        eps = DQN.compute_epsilon(dummy)
        assert abs( eps - 0.01 ) < 0.01, "ERROR: compute_epsilon at t=500000 should be .01 but g

        dummy.num_act_steps = 600000
        eps = DQN.compute_epsilon(dummy)
        assert abs( eps - 0.01 ) < 0.01, "ERROR: compute_epsilon after t=500000 should be .01 bu
        print ("Epsilon-greedy test passed!")
```

Epsilon-greedy test passed!

C3 (10 pts): Complete the code for computing the loss function in `DQN.train` in file `Algo.py` And run the cell below to verify your code decreases the loss value in one iteration.

```
In [5]: import numpy as np
        from Algo import DQN
```

```

class Nothing: pass
dummy_obs_space, dummy_act_space = Nothing(), Nothing()
dummy_obs_space.shape = [10]
dummy_act_space.n = 3

dqn = DQN(dummy_obs_space, dummy_act_space, batch_size=2)

for t in range(3):
    dqn.observe([np.random.randn(10).astype('float32')], [np.random.randint(3)],
                [(np.random.randn(10).astype('float32'), np.random.rand(), False, None)]

    b = dqn.replay.cur_batch
    loss1 = dqn.train()
    dqn.replay.cur_batch = b
    loss2 = dqn.train()

    print (loss1, '>', loss2, '?')
    assert loss2 < loss1, "DQN.train should reduce loss on the same batch"

    print ("DQN.train test passed!")

```

parameters to optimize: [('fc1.weight', torch.Size([128, 10]), True), ('fc1.bias', torch.Size([1

0.28368130326271057 > 0.28000691533088684 ?
DQN.train test passed!

P1 (10 pts): Run DQN on CartPole and plot the learning curve (i.e. averaged episodic reward against env steps). Your code should be able to achieve >150 averaged reward in 10000 iterations (20000 simulation steps) in only a few minutes. This is a good indication that the implementation is correct. It's ok that the curve is not always monotonically increasing because of randomness in training.

```

In [7]: %run Main.py \
        --niter 10000 \
        --env CartPole-v1 \
        --algo dqn \
        --nproc 2 \
        --lr 0.001 \
        --train_freq 1 \
        --train_start 100 \
        --replay_size 20000 \
        --batch_size 64 \
        --discount 0.996 \
        --target_update 1000 \
        --eps_decay 4000 \
        --print_freq 200 \
        --checkpoint_freq 20000 \

```

```
--save_dir cartpole_dqn \
--log log.txt \
--parallel_env 0
```

```
Namespace(algo='dqn', batch_size=64, checkpoint_freq=20000, discount=0.996, ent_coef=0.01, env='
observation space: Box(4,)
action space: Discrete(2)
running on device cuda
parameters to optimize: [('fc1.weight', torch.Size([128, 4]), True), ('fc1.bias', torch.Size([12
```

```
obses on reset: 2 x (4,) float32
```

iter	200	loss	0.01	n_ep	16	ep_len	21.6	ep_rew	21.59	raw_ep_rew	21.59	env_step
iter	400	loss	0.03	n_ep	33	ep_len	21.6	ep_rew	21.57	raw_ep_rew	21.57	env_step
iter	600	loss	0.01	n_ep	47	ep_len	25.8	ep_rew	25.80	raw_ep_rew	25.80	env_step
iter	800	loss	0.02	n_ep	66	ep_len	22.1	ep_rew	22.06	raw_ep_rew	22.06	env_step
iter	1000	loss	0.03	n_ep	88	ep_len	17.0	ep_rew	17.01	raw_ep_rew	17.01	env_step
iter	1200	loss	0.02	n_ep	109	ep_len	19.9	ep_rew	19.91	raw_ep_rew	19.91	env_step
iter	1400	loss	0.02	n_ep	126	ep_len	22.0	ep_rew	21.97	raw_ep_rew	21.97	env_step
iter	1600	loss	0.01	n_ep	146	ep_len	21.9	ep_rew	21.93	raw_ep_rew	21.93	env_step
iter	1800	loss	0.02	n_ep	168	ep_len	17.0	ep_rew	16.98	raw_ep_rew	16.98	env_step
iter	2000	loss	0.01	n_ep	185	ep_len	24.6	ep_rew	24.59	raw_ep_rew	24.59	env_step
iter	2200	loss	0.06	n_ep	202	ep_len	21.0	ep_rew	21.05	raw_ep_rew	21.05	env_step
iter	2400	loss	0.01	n_ep	215	ep_len	27.3	ep_rew	27.29	raw_ep_rew	27.29	env_step
iter	2600	loss	0.00	n_ep	225	ep_len	35.8	ep_rew	35.84	raw_ep_rew	35.84	env_step
iter	2800	loss	0.01	n_ep	235	ep_len	33.6	ep_rew	33.55	raw_ep_rew	33.55	env_step
iter	3000	loss	0.02	n_ep	247	ep_len	32.9	ep_rew	32.94	raw_ep_rew	32.94	env_step
iter	3200	loss	0.09	n_ep	255	ep_len	39.5	ep_rew	39.48	raw_ep_rew	39.48	env_step
iter	3400	loss	0.01	n_ep	261	ep_len	58.4	ep_rew	58.35	raw_ep_rew	58.35	env_step
iter	3600	loss	0.00	n_ep	266	ep_len	64.9	ep_rew	64.94	raw_ep_rew	64.94	env_step
iter	3800	loss	0.08	n_ep	267	ep_len	72.6	ep_rew	72.65	raw_ep_rew	72.65	env_step
iter	4000	loss	0.01	n_ep	270	ep_len	103.2	ep_rew	103.16	raw_ep_rew	103.16	env_step
iter	4200	loss	0.04	n_ep	273	ep_len	111.4	ep_rew	111.41	raw_ep_rew	111.41	env_step
iter	4400	loss	0.02	n_ep	277	ep_len	112.7	ep_rew	112.71	raw_ep_rew	112.71	env_step
iter	4600	loss	0.00	n_ep	281	ep_len	112.4	ep_rew	112.41	raw_ep_rew	112.41	env_step
iter	4800	loss	0.00	n_ep	283	ep_len	114.2	ep_rew	114.22	raw_ep_rew	114.22	env_step
iter	5000	loss	0.08	n_ep	287	ep_len	123.7	ep_rew	123.75	raw_ep_rew	123.75	env_step
iter	5200	loss	0.06	n_ep	289	ep_len	124.8	ep_rew	124.79	raw_ep_rew	124.79	env_step
iter	5400	loss	0.06	n_ep	291	ep_len	135.8	ep_rew	135.84	raw_ep_rew	135.84	env_step
iter	5600	loss	0.08	n_ep	294	ep_len	140.3	ep_rew	140.32	raw_ep_rew	140.32	env_step
iter	5800	loss	0.02	n_ep	295	ep_len	144.9	ep_rew	144.89	raw_ep_rew	144.89	env_step
iter	6000	loss	0.03	n_ep	299	ep_len	154.5	ep_rew	154.54	raw_ep_rew	154.54	env_step
iter	6200	loss	0.00	n_ep	301	ep_len	158.5	ep_rew	158.51	raw_ep_rew	158.51	env_step
iter	6400	loss	0.00	n_ep	303	ep_len	154.6	ep_rew	154.60	raw_ep_rew	154.60	env_step
iter	6600	loss	0.01	n_ep	305	ep_len	157.7	ep_rew	157.68	raw_ep_rew	157.68	env_step
iter	6800	loss	0.01	n_ep	309	ep_len	157.3	ep_rew	157.26	raw_ep_rew	157.26	env_step
iter	7000	loss	0.02	n_ep	310	ep_len	160.5	ep_rew	160.54	raw_ep_rew	160.54	env_step
iter	7200	loss	0.00	n_ep	312	ep_len	175.1	ep_rew	175.12	raw_ep_rew	175.12	env_step
iter	7400	loss	0.02	n_ep	314	ep_len	170.3	ep_rew	170.32	raw_ep_rew	170.32	env_step

```

iter    7600 |loss    0.01 |n_ep    317 |ep_len    167.8 |ep_rew    167.82 |raw_ep_rew    167.82 |env_step
iter    7800 |loss    0.04 |n_ep    320 |ep_len    166.8 |ep_rew    166.80 |raw_ep_rew    166.80 |env_step
iter    8000 |loss    0.02 |n_ep    322 |ep_len    168.0 |ep_rew    167.99 |raw_ep_rew    167.99 |env_step
iter    8200 |loss    0.09 |n_ep    324 |ep_len    169.0 |ep_rew    168.97 |raw_ep_rew    168.97 |env_step
iter    8400 |loss    0.01 |n_ep    327 |ep_len    163.0 |ep_rew    162.95 |raw_ep_rew    162.95 |env_step
iter    8600 |loss    0.08 |n_ep    330 |ep_len    166.5 |ep_rew    166.48 |raw_ep_rew    166.48 |env_step
iter    8800 |loss    0.23 |n_ep    332 |ep_len    163.3 |ep_rew    163.31 |raw_ep_rew    163.31 |env_step
iter    9000 |loss    0.15 |n_ep    334 |ep_len    162.3 |ep_rew    162.29 |raw_ep_rew    162.29 |env_step
iter    9200 |loss    0.12 |n_ep    337 |ep_len    159.2 |ep_rew    159.23 |raw_ep_rew    159.23 |env_step
iter    9400 |loss    0.16 |n_ep    339 |ep_len    163.4 |ep_rew    163.37 |raw_ep_rew    163.37 |env_step
iter    9600 |loss    0.01 |n_ep    342 |ep_len    162.7 |ep_rew    162.68 |raw_ep_rew    162.68 |env_step
iter    9800 |loss    0.18 |n_ep    344 |ep_len    161.5 |ep_rew    161.54 |raw_ep_rew    161.54 |env_step
save checkpoint to cartpole_dqn/9999.pth

```

```
In [1]: import matplotlib.pyplot as plt
```


```

def plot_curve(logfile, title=None):
    lines = open(logfile, 'r').readlines()
    lines = [l.split() for l in lines if l[:4] == 'iter']
    steps = [int(l[13]) for l in lines]
    rewards = [float(l[11]) for l in lines]
    plt.plot(steps, rewards)
    plt.xlabel('env steps'); plt.ylabel('avg episode reward'); plt.grid(True)
    if title: plt.title(title)
    plt.show()

```

The log is saved to 'cartpole_dqn/log.txt'. Let's plot the running averaged episode reward curve during training:

```
In [9]: plot_curve('cartpole_dqn/log.txt', 'CartPole DQN')
```

1.1.4 1.4 Actor-Critic Algorithm

Policy gradient methods are another class of algorithms that originated from viewing the RL problem as a mathematical optimization problem. Recall that the objective of RL is to maximize the expected cumulative reward the agent gets, namely

$$\max_{\pi} J(\pi) := \mathbb{E}_{(s_t, a_t, r_t) \sim D^{\pi}} \left[\sum_{t=0}^{\infty} \gamma^t r_t \right]$$

where D^{π} is the distribution of trajectories induced by policy π , and inside the expectation is the random variable representing the discounted cumulative reward and J is the reward (or cost) functional. Essentially, we want to optimize the policy π .

The most straightforward way is to run gradient update on the parameter θ of a *parameterized* policy π_{θ} . One such algorithm is the so-called Advantage Actor-Critic (A2C). A2C is an on-policy policy optimization type algorithm. While collecting on-policy data, we iteratively run gradient ascent:

$$\theta_{new} \leftarrow \theta_{old} + \eta \hat{\nabla}_{\theta} J(\pi_{\theta_{old}})$$

with a Monte Carlo estimate $\hat{\nabla}_{\theta} J$ of the true gradient $\nabla_{\theta} J$. The true gradient writes as (by Policy Gradient Theorem and some manipulations):

$$\nabla_{\theta} J(\pi_{\theta_{old}}) = \mathbb{E}_{(s_t, a_t, r_t) \sim D^{\pi_{\theta_{old}}}} \sum_{t=0}^{\infty} \left(\nabla_{\theta} \log \pi_{\theta_{old}}(s_t, a_t) \left(\sum_{t'=t}^{\infty} \gamma^{t'-t} r_{t'} - V^{\pi_{\theta_{old}}}(s_t) \right) \right).$$

The quantity in the inner-most parentheses $A(s_t, a_t) = Q(s_t, a_t) - V(s_t) = (\mathbb{E} \sum_{t'=t}^{\infty} \gamma^{t'-t} r_{t'}) - V(s_t)$ is called the *Advantage* function (not very rigorously speaking...). That's why it's called **Advantage** Actor-Critic. More on A2C: <https://arxiv.org/abs/1506.02438>.

And the Monte Carlo estimate of the gradient is

$$\hat{\nabla}_{\theta} J(\pi_{\theta_{old}}) = \frac{1}{NT} \sum_{i=1}^N \sum_{t=0}^T \left(\nabla_{\theta} \log \pi_{\theta_{old}}(s_t^i, a_t^i) \left(\sum_{t'=t}^T \gamma^{t'-t} r_{t'}^i - V_{\phi_{old}}(s_t^i) \right) \right)$$

where $V_{\phi_{old}}$ is introduced as a *parameterized* estimate for $V^{\pi_{\theta_{old}}}$, which can also be a neural network. So V_{ϕ} is the **critic** and π_{θ} is the **actor**. We can construct a specific loss function in pytorch that gives $\hat{\nabla}_{\theta} J$. $V_{\phi_{old}}$ is trained with SGD on a L2 loss function. It's further common practice to add an entropy bonus loss term to encourage maximum entropy solution, to facilitate exploration and avoid getting stuck in local minima. We shall clarify these loss functions in the following summarization.

Summarizing a variant of the A2C algorithm:

For many iterations repeat: 1. Collect N independent trajectories $\{(s_t^i, a_t^i, r_t^i)_{t=0}^T\}_{i=1}^N$ by running policy π_{θ} for maximum T steps; 2. Compute the loss function for the policy parameter θ :

$$L_{policy}(\theta) = \frac{1}{NT} \sum_{i=1}^N \sum_{t=0}^T \left(\log \pi_{\theta}(s_t^i, a_t^i) \left(\sum_{t'=t}^T \gamma^{t'-t} r_{t'}^i - V_{\phi}(s_t^i) \right) \right)$$

Compute the entropy term for θ :

$$L_{entropy}(\theta) = \frac{1}{NT} \sum_{i=1}^N \sum_{t=0}^T \left(- \sum_{a \in A} \pi_{\theta}(s_t^i, a) \log \pi_{\theta}(s_t^i, a) \right)$$

Compute the loss for value function parameter ϕ :

$$L_{value}(\phi) = \frac{1}{NT} \sum_{i=1}^N \sum_{t=0}^T \left(\sum_{t'=t}^T \gamma^{t'-t} r_{t'}^i - V_{\phi}(s_t^i) \right)^2$$

3. Use pytorch auto-differentiation and optimizer to do one gradient step on (θ, ϕ) with the overall loss:

$$L(\theta, \phi) = -L_{policy} - \lambda_{ent} L_{entropy} + \lambda_{val} L_{value}$$

where λ_{ent} and λ_{val} are coefficients to balances the loss terms.

C4 (10 pts): Complete the code for computing the advantage, entropy and loss function in A2C.train in file Algo.py

In []:

P2 (10 pts): Run A2C on CartPole and plot the learning curve (i.e. averaged episodic reward against training iteration). Your code should be able to achieve >150 averaged reward in 10000 iterations (40000 simulation steps) in only a few minutes. This is a good indication that the implementation is correct.

```
In [10]: %run Main.py \
        --niter 10000 \
        --env CartPole-v1 \
        --algo a2c \
        --nproc 4 \
        --lr 0.001 \
        --train_freq 16 \
        --train_start 0 \
        --batch_size 64 \
        --discount 0.996 \
        --value_coef 0.01 \
        --print_freq 200 \
        --checkpoint_freq 20000 \
        --save_dir cartpole_a2c \
        --log log.txt \
        --parallel_env 0
```

```
Namespace(algo='a2c', batch_size=64, checkpoint_freq=20000, discount=0.996, ent_coef=0.01, env='
observation space: Box(4,)
action space: Discrete(2)
running on device cuda
shared net = False, parameters to optimize: [('fc1.weight', torch.Size([128, 4])), True), ('fc1.b
```


```
obses on reset: 4 x (4,) float32
```

iter	200	loss	0.80	n_ep	25	ep_len	26.5	ep_rew	26.49	raw_ep_rew	26.49	env_step
iter	400	loss	0.87	n_ep	54	ep_len	26.7	ep_rew	26.68	raw_ep_rew	26.68	env_step
iter	600	loss	1.03	n_ep	76	ep_len	36.4	ep_rew	36.36	raw_ep_rew	36.36	env_step
iter	800	loss	0.72	n_ep	98	ep_len	41.5	ep_rew	41.48	raw_ep_rew	41.48	env_step
iter	1000	loss	0.85	n_ep	122	ep_len	32.2	ep_rew	32.18	raw_ep_rew	32.18	env_step
iter	1200	loss	0.82	n_ep	137	ep_len	49.3	ep_rew	49.29	raw_ep_rew	49.29	env_step
iter	1400	loss	0.75	n_ep	156	ep_len	39.2	ep_rew	39.23	raw_ep_rew	39.23	env_step
iter	1600	loss	0.86	n_ep	173	ep_len	41.4	ep_rew	41.40	raw_ep_rew	41.40	env_step
iter	1800	loss	0.95	n_ep	191	ep_len	42.6	ep_rew	42.61	raw_ep_rew	42.61	env_step
iter	2000	loss	0.51	n_ep	201	ep_len	73.8	ep_rew	73.81	raw_ep_rew	73.81	env_step
iter	2200	loss	0.66	n_ep	217	ep_len	55.2	ep_rew	55.19	raw_ep_rew	55.19	env_step
iter	2400	loss	0.81	n_ep	232	ep_len	56.1	ep_rew	56.10	raw_ep_rew	56.10	env_step
iter	2600	loss	0.92	n_ep	243	ep_len	47.9	ep_rew	47.86	raw_ep_rew	47.86	env_step
iter	2800	loss	1.03	n_ep	254	ep_len	70.7	ep_rew	70.71	raw_ep_rew	70.71	env_step
iter	3000	loss	1.00	n_ep	262	ep_len	80.1	ep_rew	80.12	raw_ep_rew	80.12	env_step
iter	3200	loss	0.57	n_ep	275	ep_len	75.1	ep_rew	75.05	raw_ep_rew	75.05	env_step
iter	3400	loss	1.00	n_ep	282	ep_len	83.4	ep_rew	83.38	raw_ep_rew	83.38	env_step
iter	3600	loss	0.96	n_ep	288	ep_len	110.2	ep_rew	110.25	raw_ep_rew	110.25	env_step
iter	3800	loss	0.67	n_ep	294	ep_len	113.5	ep_rew	113.49	raw_ep_rew	113.49	env_step

iter	4000	loss	0.83	n_ep	297	ep_len	129.2	ep_rew	129.21	raw_ep_rew	129.21	env_step
iter	4200	loss	0.97	n_ep	304	ep_len	154.1	ep_rew	154.09	raw_ep_rew	154.09	env_step
iter	4400	loss	0.28	n_ep	311	ep_len	137.0	ep_rew	137.03	raw_ep_rew	137.03	env_step
iter	4600	loss	0.92	n_ep	317	ep_len	125.9	ep_rew	125.93	raw_ep_rew	125.93	env_step
iter	4800	loss	0.86	n_ep	325	ep_len	119.8	ep_rew	119.81	raw_ep_rew	119.81	env_step
iter	5000	loss	0.83	n_ep	332	ep_len	105.7	ep_rew	105.68	raw_ep_rew	105.68	env_step
iter	5200	loss	0.29	n_ep	337	ep_len	135.9	ep_rew	135.90	raw_ep_rew	135.90	env_step
iter	5400	loss	0.33	n_ep	345	ep_len	117.9	ep_rew	117.90	raw_ep_rew	117.90	env_step
iter	5600	loss	1.00	n_ep	350	ep_len	136.1	ep_rew	136.14	raw_ep_rew	136.14	env_step
iter	5800	loss	1.00	n_ep	355	ep_len	139.9	ep_rew	139.90	raw_ep_rew	139.90	env_step
iter	6000	loss	0.18	n_ep	361	ep_len	130.7	ep_rew	130.68	raw_ep_rew	130.68	env_step
iter	6200	loss	0.72	n_ep	368	ep_len	130.4	ep_rew	130.41	raw_ep_rew	130.41	env_step
iter	6400	loss	0.49	n_ep	375	ep_len	123.4	ep_rew	123.42	raw_ep_rew	123.42	env_step
iter	6600	loss	0.86	n_ep	379	ep_len	141.5	ep_rew	141.55	raw_ep_rew	141.55	env_step
iter	6800	loss	0.82	n_ep	383	ep_len	152.9	ep_rew	152.87	raw_ep_rew	152.87	env_step
iter	7000	loss	0.31	n_ep	389	ep_len	141.1	ep_rew	141.10	raw_ep_rew	141.10	env_step
iter	7200	loss	0.70	n_ep	393	ep_len	159.0	ep_rew	159.01	raw_ep_rew	159.01	env_step
iter	7400	loss	1.11	n_ep	397	ep_len	162.5	ep_rew	162.48	raw_ep_rew	162.48	env_step
iter	7600	loss	0.16	n_ep	401	ep_len	193.9	ep_rew	193.87	raw_ep_rew	193.87	env_step
iter	7800	loss	0.66	n_ep	404	ep_len	203.2	ep_rew	203.21	raw_ep_rew	203.21	env_step
iter	8000	loss	0.28	n_ep	409	ep_len	198.1	ep_rew	198.14	raw_ep_rew	198.14	env_step
iter	8200	loss	0.99	n_ep	410	ep_len	207.9	ep_rew	207.93	raw_ep_rew	207.93	env_step
iter	8400	loss	0.87	n_ep	416	ep_len	202.1	ep_rew	202.08	raw_ep_rew	202.08	env_step
iter	8600	loss	-0.08	n_ep	420	ep_len	215.6	ep_rew	215.61	raw_ep_rew	215.61	env_step
iter	8800	loss	0.79	n_ep	425	ep_len	201.7	ep_rew	201.66	raw_ep_rew	201.66	env_step
iter	9000	loss	0.74	n_ep	431	ep_len	156.5	ep_rew	156.51	raw_ep_rew	156.51	env_step
iter	9200	loss	0.30	n_ep	434	ep_len	155.4	ep_rew	155.39	raw_ep_rew	155.39	env_step
iter	9400	loss	0.61	n_ep	442	ep_len	151.9	ep_rew	151.88	raw_ep_rew	151.88	env_step
iter	9600	loss	0.64	n_ep	445	ep_len	143.6	ep_rew	143.59	raw_ep_rew	143.59	env_step
iter	9800	loss	0.91	n_ep	451	ep_len	163.1	ep_rew	163.09	raw_ep_rew	163.09	env_step

save checkpoint to cartpole_a2c/9999.pth

In [11]: plot_curve('cartpole_a2c/log.txt', 'CartPole A2C')



Now let's play a little bit with the trained agent. The neural net parameters are saved to the `cartpole_dqn` and `cartpole_a2c` folders. The cell below will open a window showing one episode play.

```
In [13]: import time
import gym
import Algo
env = gym.make('CartPole-v1')
agent = Algo.ActorCritic(env.observation_space, env.action_space)
agent.load('cartpole_a2c/9999.pth')
state = env.reset()
for _ in range(120):
    env.render()
    state, reward, done, _ = env.step(agent.act([state])[0])
    if done: break
    time.sleep(0.1)
env.close()
```

```
shared net = False, parameters to optimize: [('fc1.weight', torch.Size([128, 4]), True), ('fc1.b
```

1.2 Part II: Solve the Atari Breakout game

In this part, you'll train your agent to play Breakout with the BlueWaters cluster. I have provided the job scripts for you. Please upload your `Algo.py` and `Model.py` completed in **Part I** to your BlueWaters folder. And submit the following two jobs respectively:

```
qsub run_dqn.pbs
qsub run_a2c.pbs
```

The jobs are set to run for at most **14 hours**. **Please start early!!** You might be able to reach the desired score (≥ 200 reward) before 14 hours - You can stop the training early if you wish. Then please collect the resulting `breakout_dqn/log.txt` and `breakout_a2c/log.txt` files into the same folder as this Jupyter notebook's. Rename them as `log_breakout_dqn.txt` and `log_breakout_a2c.txt`.

BTW, there's an Atari PC simulator: <https://stella-emu.github.io/> I spent a lot of time playing them...


C5 (10 pts): Complete the code for the CNN with 3 conv layers and 3 fc layers in class SimpleCNN in file Model.py And verify the output shape with the cell below.

```
In [6]: ## Test code
        from Model import SimpleCNN
        import torch
        net = SimpleCNN()
        x = torch.randn(2, 4, 84, 84)
        y = net(x)
        assert y.shape == (2, 4), "ERROR: network output has the wrong shape!"
        print("CNN output shape test passed!")
```

CNN output shape test passed!


P3 (10 pts): Run the following cell to generate a DQN learning curve. The *maximum* average episodic reward on this curve should be larger than 200 for full credit. (It's ok if the final reward is not as high.) The typical value is around 300. You get 70% credit if $100 \leq \text{average episodic reward} < 200$, 50% credit if $50 \leq \text{average episodic reward} < 100$.

```
In [14]: plot_curve('log_breakout_dqn.txt', 'Breakout DQN')
```



P4 (10 pts): Run the following cell to generate an A2C learning curve. The *maximum* average episodic reward on this curve should be larger than 150 for full credit. (It's ok if the final reward is not as high.) The typical value is around 250. You get 70% credit if $50 \leq \text{average episodic reward} < 150$, and 50% credit if $20 \leq \text{average episodic reward} < 50$.

```
In [2]: plot_curve('log_breakout_a2c.txt', 'Breakout A2C')
```




P5 (10 pts): Collect and visualize some game frames by running the script `Draw.py` on BlueWaters.

- (1) module load python/2.0.0 and run `Draw.py` on BlueWaters (it's ok to run this locally, no need to start a job).
- (2) Download the result `breakout_imgs` folder from BlueWaters to the folder containing this Jupyter notebook, and run the following cell. You should see some animation of the game.

```
In [16]: import os
         imgs = sorted(os.listdir('breakout_imgs'))
         #imgs = [plt.imread('breakout_imgs/' + img) for img in imgs]

         %matplotlib inline
         import matplotlib.pyplot as plt
         from IPython import display
         pimg = None
         for img in imgs:
             img = plt.imread('breakout_imgs/' + img)
             if pimg:
                 pimg.set_data(img)
             else:
                 pimg = plt.imshow(img)
             display.display(plt.gcf())
             display.clear_output(wait=True)
```

1.3 Part III: Questions (10 pts)

These are open-ended questions. The purpose is to encourage you to think (a bit) more deeply about these algorithms. You get full points as long as you write a few sentences that make sense and show some thinking.

Q1 (2 pts): Why would people want to do function approximation rather than using tabular algorithm (on discretized S, A spaces if necessary)? Bringing function approximation has caused numerous problems theoretically (e.g. not guaranteed to converge), so it seems not worth it... Your answer: I don't know. People enjoy "neuralizing" things I guess..

Q2 (2 pts): Q-Learning seems good... it's theoretically sound (at least seems to be), the performance is also good. Why would many people actually prefer policy gradient type algorithms in some practical problems? Your answer: I don't know. I like Q learning. The name is cute. Anyone watch StarTrek?

Q3 (2 pts): Does the policy gradient algorithm (A2C) we implemented here extend to continuous action space? How would you do that? Hint: What is a reasonable distribution assumption for policy $\pi_{\theta}(a|s)$ if a lives in continuous space? Your answer: I don't know. Maybe normalizing flow?? OK, people really do this..(arXiv:1905.06893) Hot area + hot area

Q4 (2 pts): The policy gradient algorithm (A2C) we implemented uses on-policy data. Can you think of a way to extend it to utilize off-policy data? **Hint: Importance sampling, needs some approximation though** Your answer: I don't know. Do random math tricks or pray?

Q5 (2 pts): How to compare different RL algorithms? When can I say one algorithm is better than the other? **Hint: This question is quite open. Think about speed, complexity, tasks, etc.** Your answer: I don't know. Just pick one you like, they're equally bad..